



Intelligently Evolving Gradient Descent

John Vahedi

Abstract

- Intelligent form of the gradient descent algorithm
- "Semi-stochastic," informed by short-term history
- 3 evolutionary methods
- Compared to plain SGD
- Linear and non-linear examples
- For stable but complex neural & transformer architectures

1. Background
2. Algorithms
3. Method
4. Results & Analysis
5. Conclusion

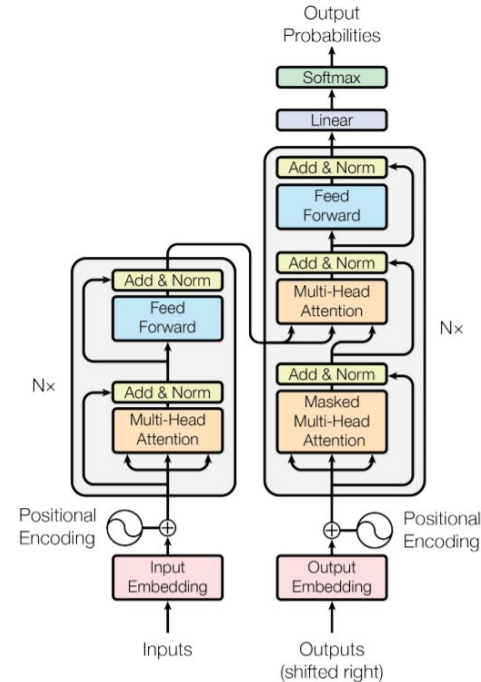


Figure 1: The Transformer - model architecture.



1.

Background

Introduction



Cost Function


Inspired Example:

❖ Business

- Reduce cost
- Two dependent choices ($n=2$)
- Current choices set

Note: Our business will be something different
And 'n' may be > 2

What is Gradient Descent?



Gradient Descent Step

Mathematically:

$$w_{i+1} = w_i \ominus \eta_i \cdot \nabla_w f(w_i) + \epsilon, \quad \eta_i = \frac{\eta_0}{i^p}, \quad p \in [0, 1]$$

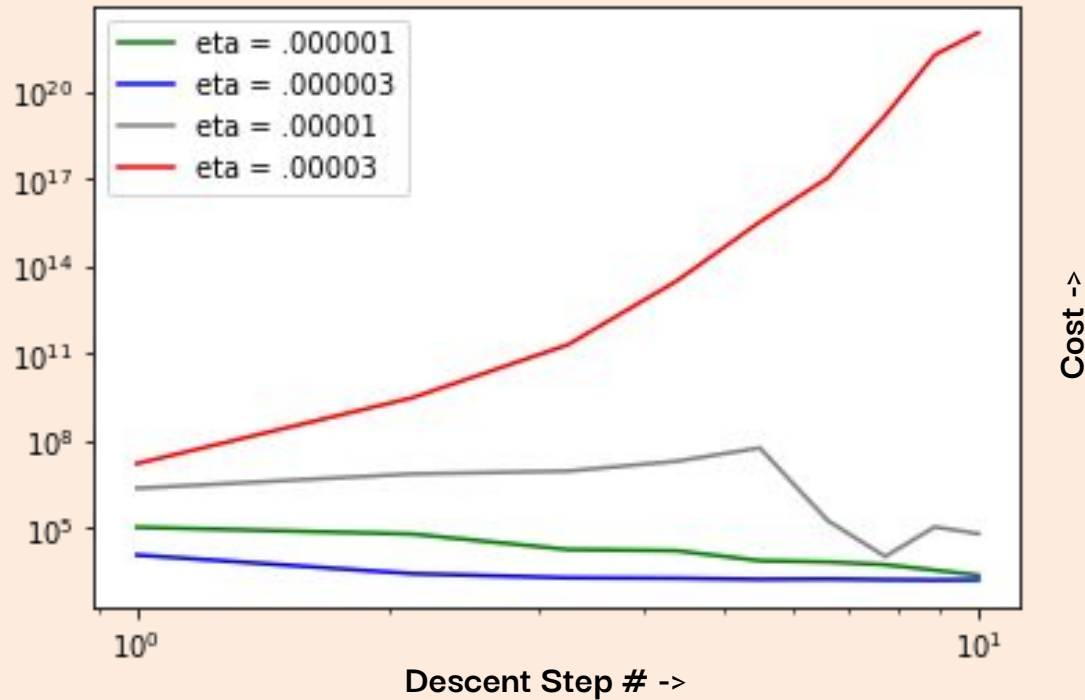
New position

Current
position

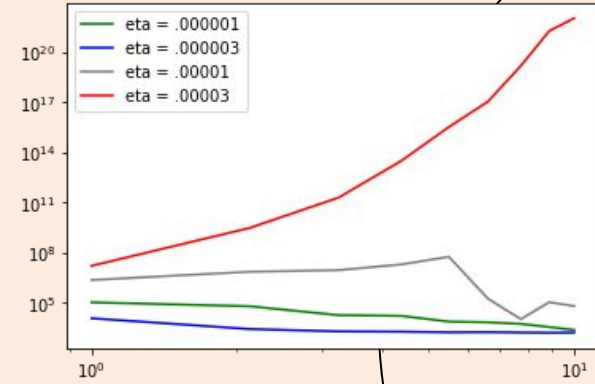

Gradient
Direction

Step Size

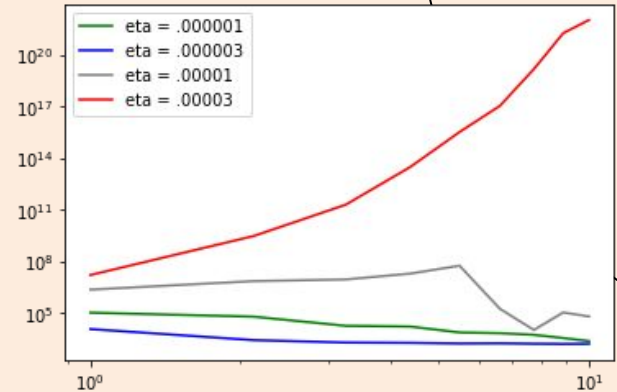
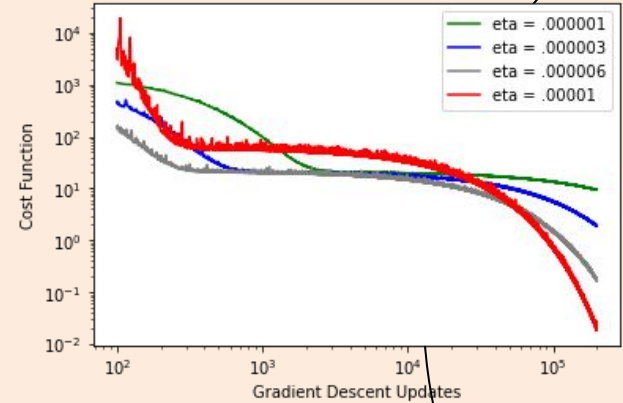

Step Size (3 Faux Pas)



Step Size (3 Faux Pas)



Step Size (3 Faux Pas)



Why Stochastic?

Definition:

- “Randomness”
- Avoiding some directions in Gradient randomly
 - Picks **subset** of gradient

Reason/strategy:

- Redundant: Data or Features
- Avoid Local Minima
- Memory Advantage

Regardless:

- Eventually reach all data (Downside Low, Upside High)



2.

Algorithms

Historical Gradient Descent



Momentum

Remembers old motion;
piecemeals out same
motion later



Noise

Explicit noise to added
gradient; helps
initialization



Adaptive

Adjust step size based
on gradient or historical
frequency



Implicit

Prescient methods;
Add complexity but
more stable



Too much stochastic....?

- ❖ Stochasticity has Benefits
 - Discussed
- ❖ Total Stochastic Choice
 - Leaves information on the table

Novel Gradient Descent

Semi-Stochastic Gradient Descent (SSGD)

- *Remembers last step*
- *Gradient as fitness*
- *Percent fittest, rolls over*

Note: Requires sorting

Novel Gradient Descent


Algorithm 1 SSGD

```
1: Import data
2: Initialize subset of data, size  $K$  - call this subset  $L_k$ 
3: Initialize  $a$  at initial guess  $a_0$ 
4: for  $i = 1, 2, \dots, N$  do
5:   Set  $\nabla_a \Phi = \text{zeros}(\text{input space dimensionality})$ 
6:   for  $point = 1, 2, \dots, K$  do
7:     Assess gradient at point  $L[\text{point}]$  and add to  $\nabla_a \Phi$ 
8:   end for
9:   Update  $a = a - (\eta/i^p) * (\nabla_a \Phi / K)$ 
10:  Selection on  $L_k$ , return  $L_{k,selected}$  holding highest-contributing points
11:  Repopulate  $L_{k,selected}$  with random points from dataset, set repopulated
    array equal to  $L_k$ 
12: end for
```

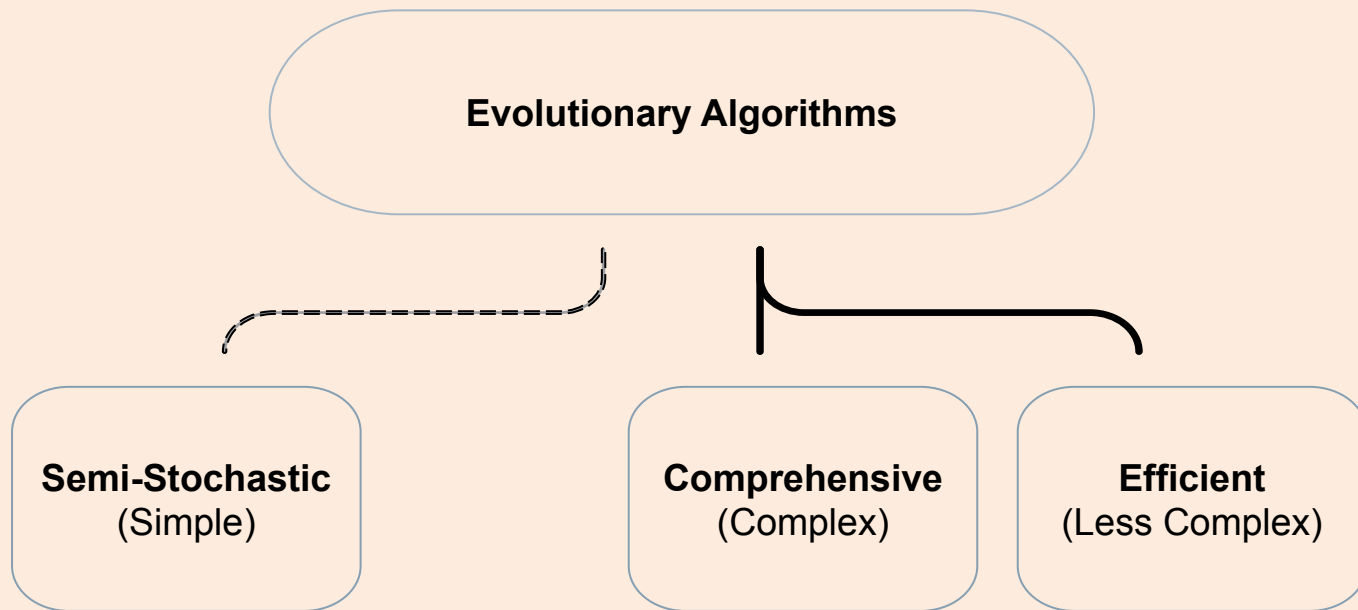
“

“Those who cannot remember the past are
condemned to repeat it.”

- **George Santayana**



Evolutionary Gradient Descent

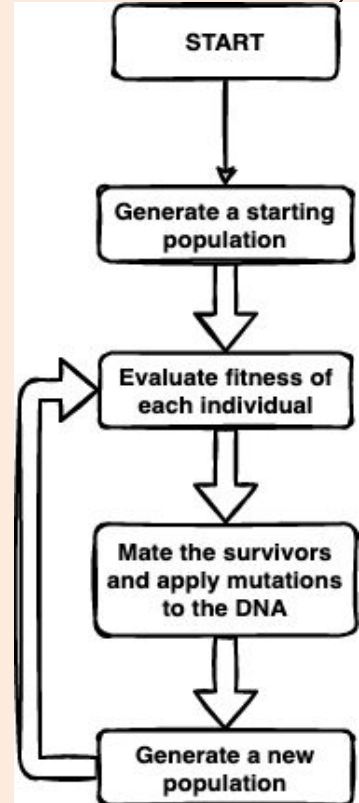


What is Evolutionary Algorithm?

- “Population” of solutions running in parallel
- Inspired by biological systems

Relevant mechanisms:

- *Fitness*
- *Breeding*
 - *Recombine*
 - *Mutation*



Evolutionary Gradient Descent

- A. Initialize P random subsets
- B. Calculate accumulated individual **Gradient**
- C. Update **Parameters** individually and **Error**
- D. Select *fittest* percentage for **Breeding**

Fitness

3 Factors:

$$fitness = |\nabla_w \Phi^*| * \frac{1}{factor_{age}} * \frac{1}{(1 + |error|)}$$

$$factor_{age} = \begin{cases} 1 & , age < C \\ (age - C) + 1, & otherwise \end{cases}$$

Evolutionary Gradient Descent

2. *Comprehensive Stochastic Gradient Descent (CEvSGD)*
 - Tracks individual's parameters and stored
 - ◆ Returns parameter with lowest error
3. *Efficient Stochastic Gradient Descent (EEvSGD)*
 - Does not track individual's parameters
 - ◆ Global parameter updated with averaged gradient

Evolutionary Gradient Descent

Algorithm 2 EvSGD Variant 1: CEvSGD (Comprehensive EvSGD)

```
1: Import data
2: Initialize  $P$  subsets of data, each of size  $K$  - call this array  $L$ 
3: Initialize array  $A$  of size  $P$ , entries equal to initial guess  $a_0$ 
4: for  $i = 1, 2, \dots, N$  do
5:   for  $member = 1, 2, \dots, P$  do
6:     Set  $L_k = L[member]$ 
7:     Set  $\nabla_a \Phi = \text{zeros}(\text{input space dimensionality})$ 
8:     for  $point = 1, 2, \dots, K$  do
9:       Assess gradient at point  $L_k[point]$  and add to  $\nabla_a \Phi[point]$ 
10:    end for
11:     $a = A[member]$ 
12:    Update  $a = a - (\eta/i^p) * (\nabla_a \Phi[member]/K)$  and correspondingly update  $A$ 
13:  end for
14:  Selection on  $L$ , return  $L_{selected}$  which holds top performing subsets
15:  Run breeding on  $L_{selected}$ , return  $L_{new}$ 
16:  Update  $A$ , with new members of the population initialized at  $A[parent]$ 
17: end for
```

Algorithm 3 EvSGD Variant 2: EEvSGD (Efficient EvSGD)


```
1: Import data
2: Initialize an array  $L$  of  $P$  subsets of data, each of size  $K$  -  $L$  corresponds to our population
3: Initialize  $a$  equal to initial guess  $a_0$ 
4: Initialize an array  $\nabla_a \Phi$  of size  $P$ , each entry is  $\text{zeros}(P)$ 
5: for  $i = 1, 2, \dots, N$  do
6:   for  $member = 1, 2, \dots, P$  do
7:     Set  $L_k = L[member]$ 
8:     Set  $\nabla_a \Phi[member] = \text{zeros}(P)$ 
9:     for  $point = 1, 2, \dots, K$  do
10:      Assess gradient at  $L_k[point]$  and add to  $\nabla_a \Phi[member]$ 
11:    end for
12:  end for
13:  Selection on  $L$ , return  $L_{selected}$ , which holds top-performing subsets
14:   $\nabla_{a,avg} \Phi = \text{average}(\nabla_a \Phi)$ 
15:  Update  $a = a - (\eta/i^p) * (\nabla_{a,avg} \Phi/K)$ 
16:  Run breeding on  $L_{selected}$ , return  $L_{new}$ 
17: end for
```

Evolutionary Gradient Descent

1. Semi-Stochastic Gradient Descent (SSGD)
2. Comprehensive Stochastic Gradient Descent (CEvSGD)
3. Efficient Stochastic Gradient Descent (EEvSGD)

Other “similar” algos

- “Evolutionary stochastic gradient descent for optimization of deep neural networks” by Xiaodong Cui et al.
- Momentum based SGD



3.

Method

Imported Packages



MatPlotLib

Generate graphs of plots



Numpy

Used for parallel
computation, and random
value generator

Regression Problem

$$F(a, x) = a_1 \cdot f_1(x) + a_2 \cdot f_2(x) + \dots a_N \cdot f_N(x) = \sum_{n=1}^N a_n \cdot f_n(x) \quad (3)$$

$$\nabla_a F(a, x) = [f_1(x), f_2(x), \dots, f_N(x)]^T \quad (4)$$

$$\Phi(a) = \frac{1}{M} \sum_{j=1}^M \phi_j(a). \quad \phi_j(a) = (y_j - F(a, x_j))^2.$$

$$\nabla_a \Phi(a) = \frac{1}{M} \sum_{j=1}^M \nabla_a \phi_j(a) = \frac{1}{M} \sum_{j=1}^M 2 \cdot (y_j - F(a, x_j)) \cdot \nabla_a F(a, x_j) \quad (8)$$

Regression Problem

$$\Phi^*(a) = \frac{1}{\sum_{j=1}^M w_j} \sum_{j=1}^M w_j \phi_j(a) = \frac{1}{K} \sum_{j=1}^M w_j \phi_j(a)$$

$$\nabla_a \Phi(a) \approx \nabla_a \Phi^*(a).$$

$$\nabla_a \Phi(a) = \nabla_a \Phi^*(a) + err, \quad err \sim N(0, \sigma) \Rightarrow$$

$$E[\nabla_a \Phi(a)] = E[\nabla_a \Phi^*(a)]$$

$$a_{i+1} = a_i - \eta_i \cdot \nabla_a \Phi^*(a_i) \tag{10}$$

Created Packages



Utils



EA_Utils



Generator



Dependencies

none

Parameter Tuning



η (eta)

Initial step size



p

Power of series used in
division to slow down
step sizes over iteration



other

%Retained, %Mutation,
%Inheritance, Age Cutoff


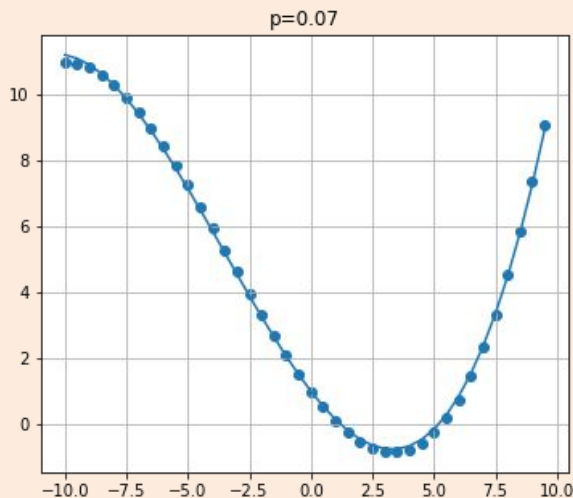
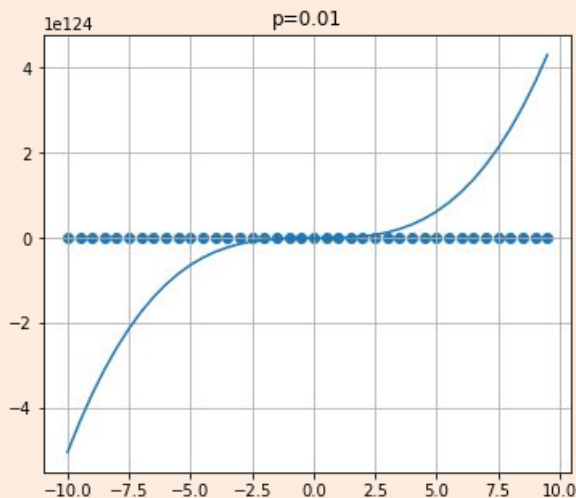
Parameter Tuning



p

Power of series used in
division to slow down
step sizes over iteration

Effect of different p values for a fixed $\eta=0.000001$ on SSGD



Parameter Tuning

Algorithm	Optimal p	Optimal η_0
1: SSGD	0.07	9×10^{-6}
2: EEvSGD	0	1.5×10^{-7}
3: CEvSGD	0.07	8.1×10^{-6}

Search done: p in $[0, 0.5]$ and η in $[1 \times 10^{-7}, 5 \times 10^{-5}]$


Normalization

Normalized Discovered Version:

$$\begin{aligned} y &= a_0^* + a_1^*(X)_s + a_2^*(X^2)_s + a_3^*(X^3)_s \\ &= a_0^* + a_1^* \frac{(X - \mu_1)}{\sigma_1} + a_2^* \frac{(X^2 - \mu_2)}{\sigma_2} + a_3^* \frac{(X^3 - \mu_3)}{\sigma_3} \\ &= \left(a_0^* - \sum_{i=0}^3 \frac{a_i^* \mu_i}{\sigma_i} \right) + \frac{a_1^*}{\sigma_1} X + \frac{a_2^*}{\sigma_2} X^2 + \frac{a_3^*}{\sigma_3} X^3 \end{aligned}$$

$$a_0 = a_0^* - \sum_{i=0}^3 \frac{a_i^* \mu_i}{\sigma_i}, \quad a_i = \frac{a_i^*}{\sigma_i}$$

- Helps Stability



4.

Results & Analysis




Figure 4: Models regressed over dataset at 30k steps (above)



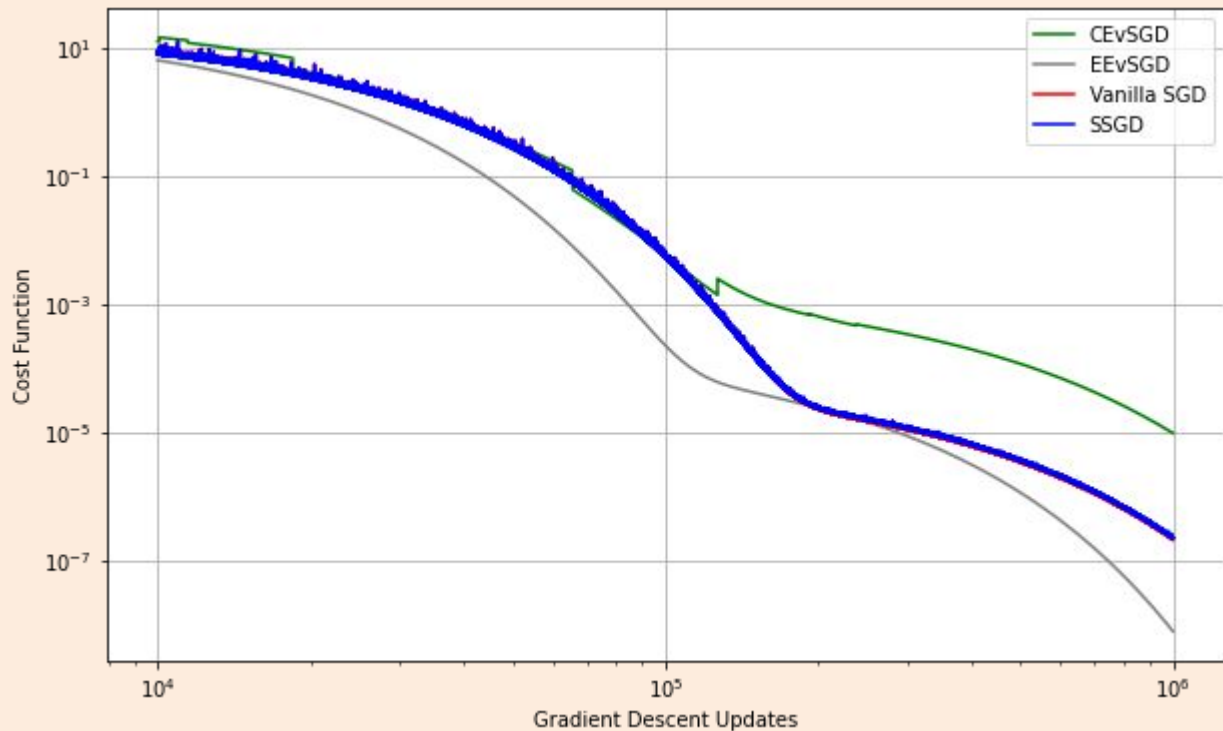


Figure 5: Models regressed over dataset at 60k steps (above)

5 Avg. Runs




5 Avg. Runs



- Higher Accuracy



Time Complexity (n runs)





~100X
Accuracy

High Accuracy:long VS fast:less accuracy

Stability



- ❖ CEvSGD ~ 50%
- ❖ EEvSGD ~ 80%
- ❖ SGD/ SSGD > 95%



5.

Conclusion & Future

Issues

- ❖ Normalization
- ❖ Stability Considerations
- ❖

Success

- ❖ CEvSGD most robust
 - Even if unstable
 - Restart
- ❖ EEvSGD most efficient and accurate
 - 100x accurate
- ❖ SSGD performs better per iteration

Speedups

- ❖ Use different norm
- ❖ Parallelization
- ❖ Find and remove redundant calculations
- ❖ Better sort or remove if possible

Future

- ❖ Non-Linear
- ❖ Parallelization
- ❖ More types of Parameter Tuning
- ❖ Fitness Criteria

References

- [1] William M Spears. “Crossover or mutation?” In: *Foundations of genetic algorithms*. Vol. 2. Elsevier, 1993, pp. 221–237.
- [2] David E Goldberg. “Genetic and evolutionary algorithms come of age”. In: *Communications of the ACM* 37.3 (1994), pp. 113–120.
- [3] Panos Toulis and Edoardo M Airolidi. “Implicit stochastic gradient descent for principled estimation with large datasets”. In: *ArXiv e-prints* (2014).
- [4] Xiaodong Cui et al. “Evolutionary stochastic gradient descent for optimization of deep neural networks”. In: *arXiv preprint arXiv:1810.06773* (2018).

**Thank
you.**

